



**Kubernetes**  
**EDGE**



**Kubernetes**  
**EDGE**

# Taming Unbounded Resources with Node Feature Discovery

*Mark Abrams*

# What is Node Feature Discovery (NFD)



Kubernetes  
EDGE

- NFD is simply a fancy labeler
- K8s node labels



# Why do I need a labeler



Kubernetes  
EDGE

- Mapping workloads to appropriate resources

Resources



Workloads

# Why do I need a labeler



**Kubernetes**  
**EDGE**

- Because we are delivering cattle not pets
- Volume - a large number of devices
  - Efficiency
  - Accuracy
- Resources at the Edge
  - Are not just CPU and RAM
  - Are Unbounded

# Unbounded Resources?



**Kubernetes**  
**EDGE**

- K8s is designed to schedule workloads against resources
- In the data center (first class features of k8s)
  - CPU
  - RAM
  - GPU
- At the edge (not schedulable by default)
  - Unlimited number of I/O devices are possible
    - Sensors
    - Actuators

# w/o NFD - Childs Play



Kubernetes  
EDGE

- Do you have any 3's



- Go fish

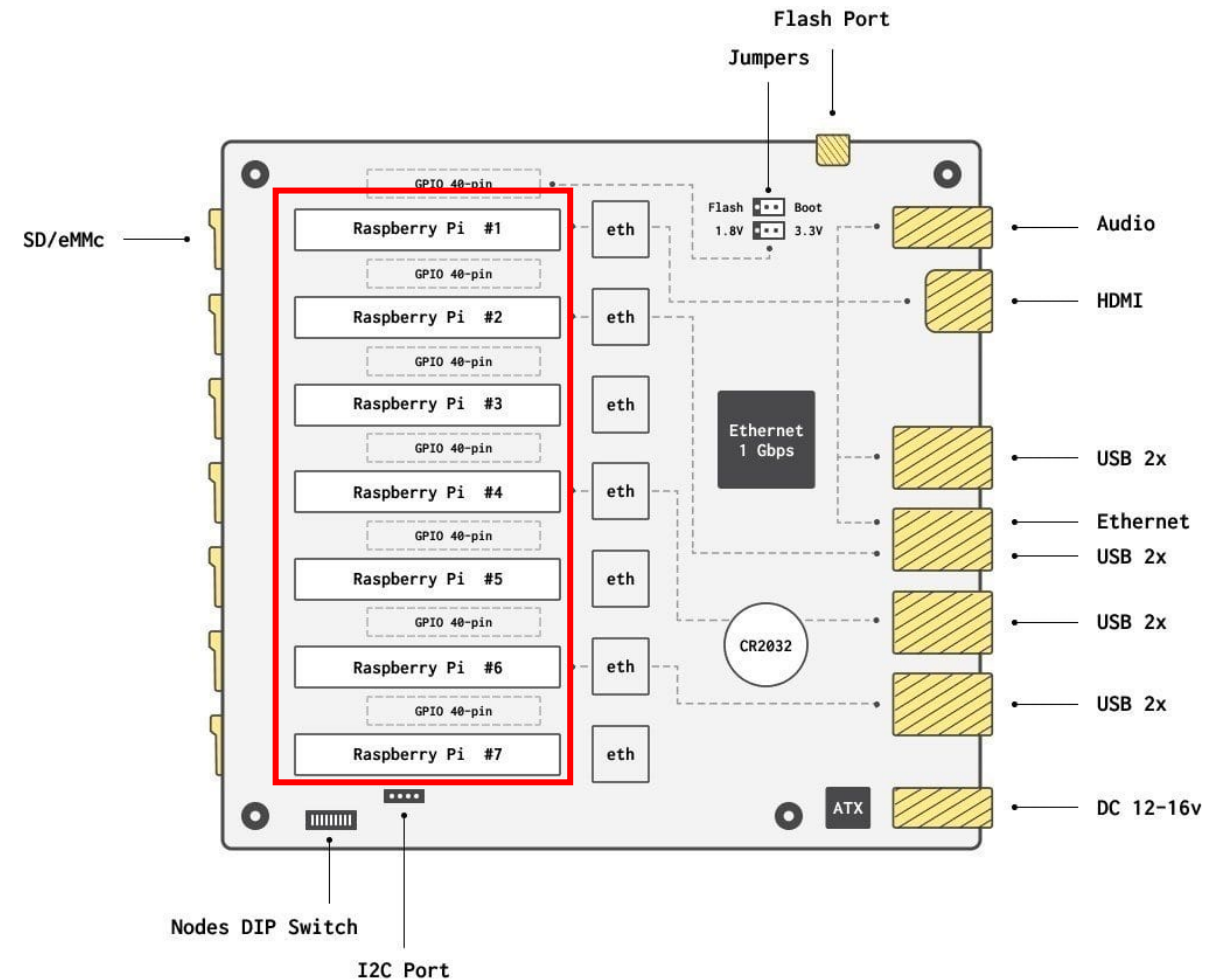


# Demo Infrastructure - turingpi v1



Kubernetes  
EDGE

- 7 RaspberryPi CM3+ modules
  - 1GB RAM
  - 4 CPU
- 8 USB ports
  - Map to only 4 devices
- I2C interface
  - Maps to all devices
  - Communication between devices
- GPIO
  - 40 pins for each device





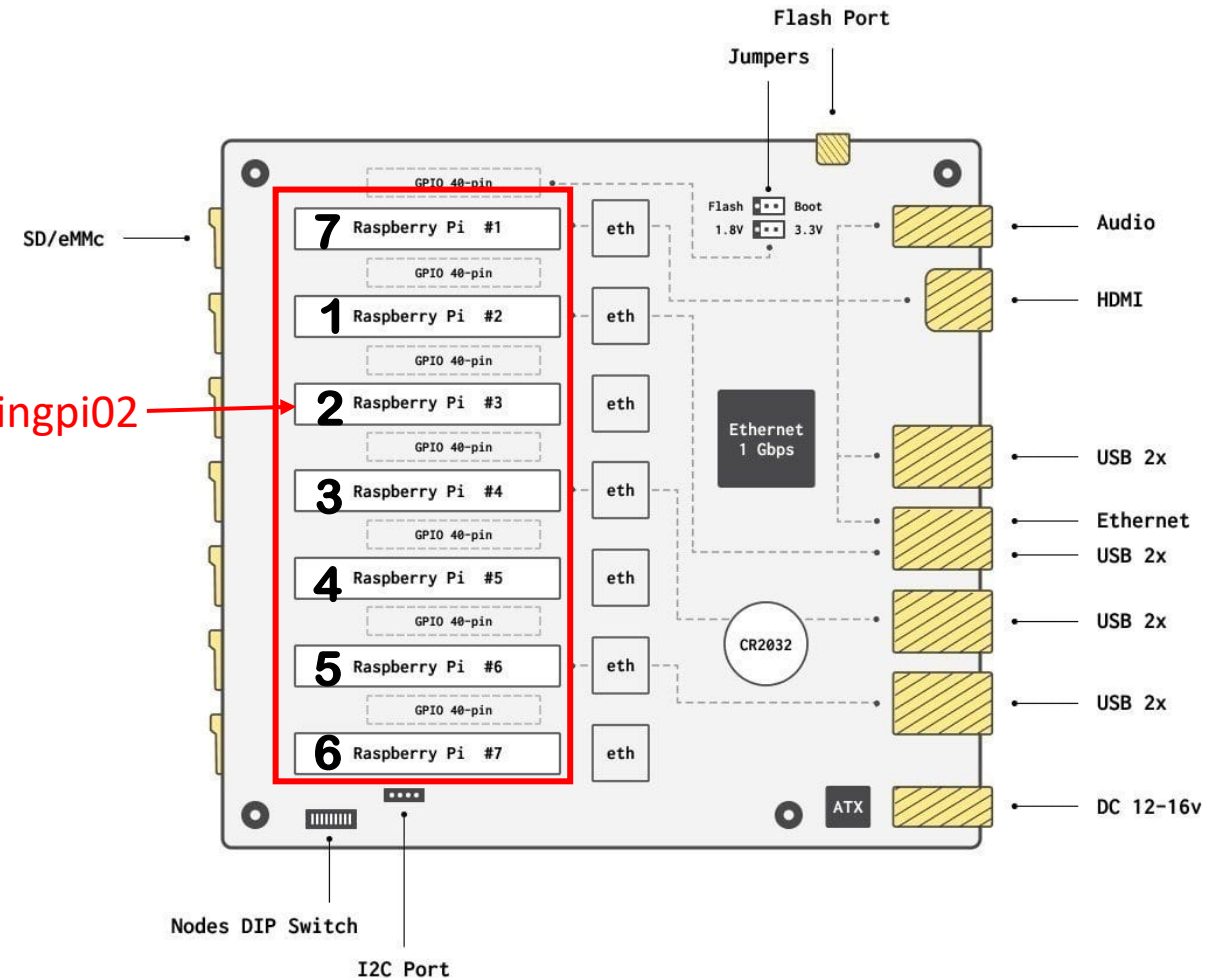
# Demo Infrastructure - turingpi v1



Kubernetes  
EDGE

- 7 RaspberryPi CM3+ modules
  - 1GB RAM
  - 4 CPU
- 8 USB ports
  - Map to only 4 devices
- I2C interface
  - Maps to all devices
  - Communication between devices
- GPIO
  - 40 pins for each device

Hostname: turingpi02

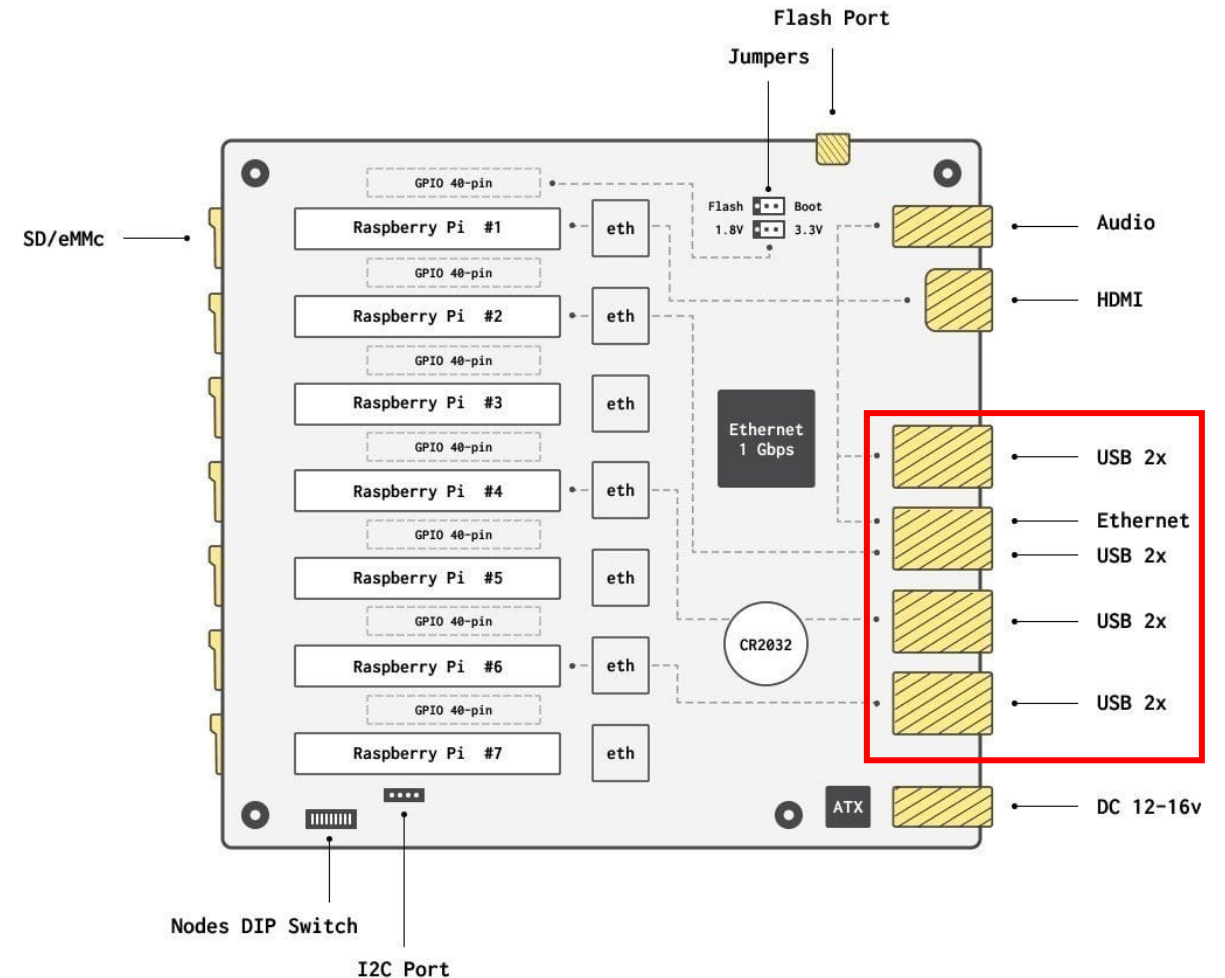


# Demo Infrastructure - turingpi v1



Kubernetes  
EDGE

- 7 RaspberryPi CM3+ modules
  - 1GB RAM
  - 4 CPU
- 8 USB ports
  - Map to only 4 devices
- I2C interface
  - Maps to all devices
  - Communication between devices
- GPIO
  - 40 pins for each device

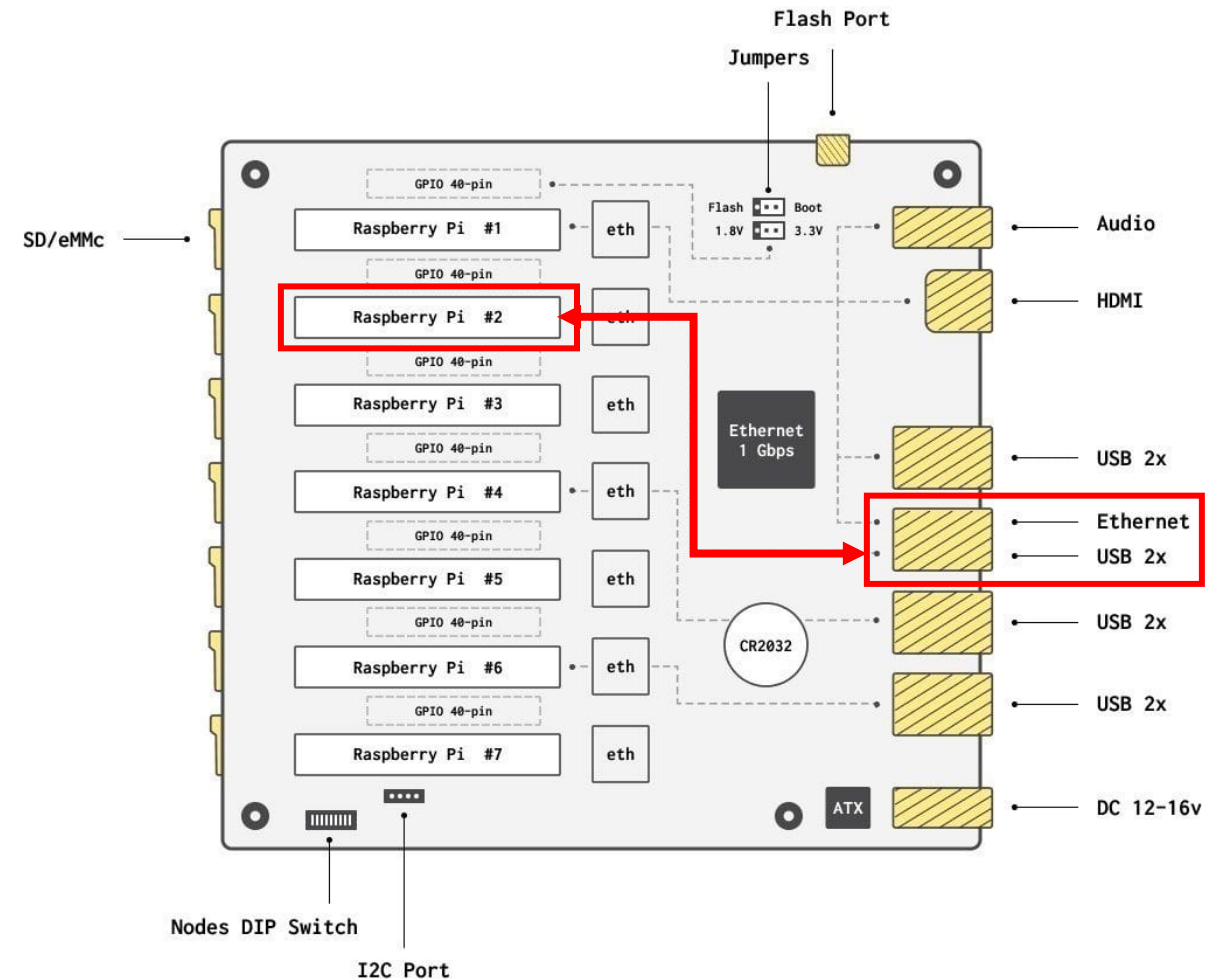


# Demo Infrastructure - turingpi v1



Kubernetes  
EDGE

- 7 RaspberryPi CM3+ modules
  - 1GB RAM
  - 4 CPU
- 8 USB ports
  - Map to only 4 devices
- I2C interface
  - Maps to all devices
  - Communication between devices
- GPIO
  - 40 pins for each device

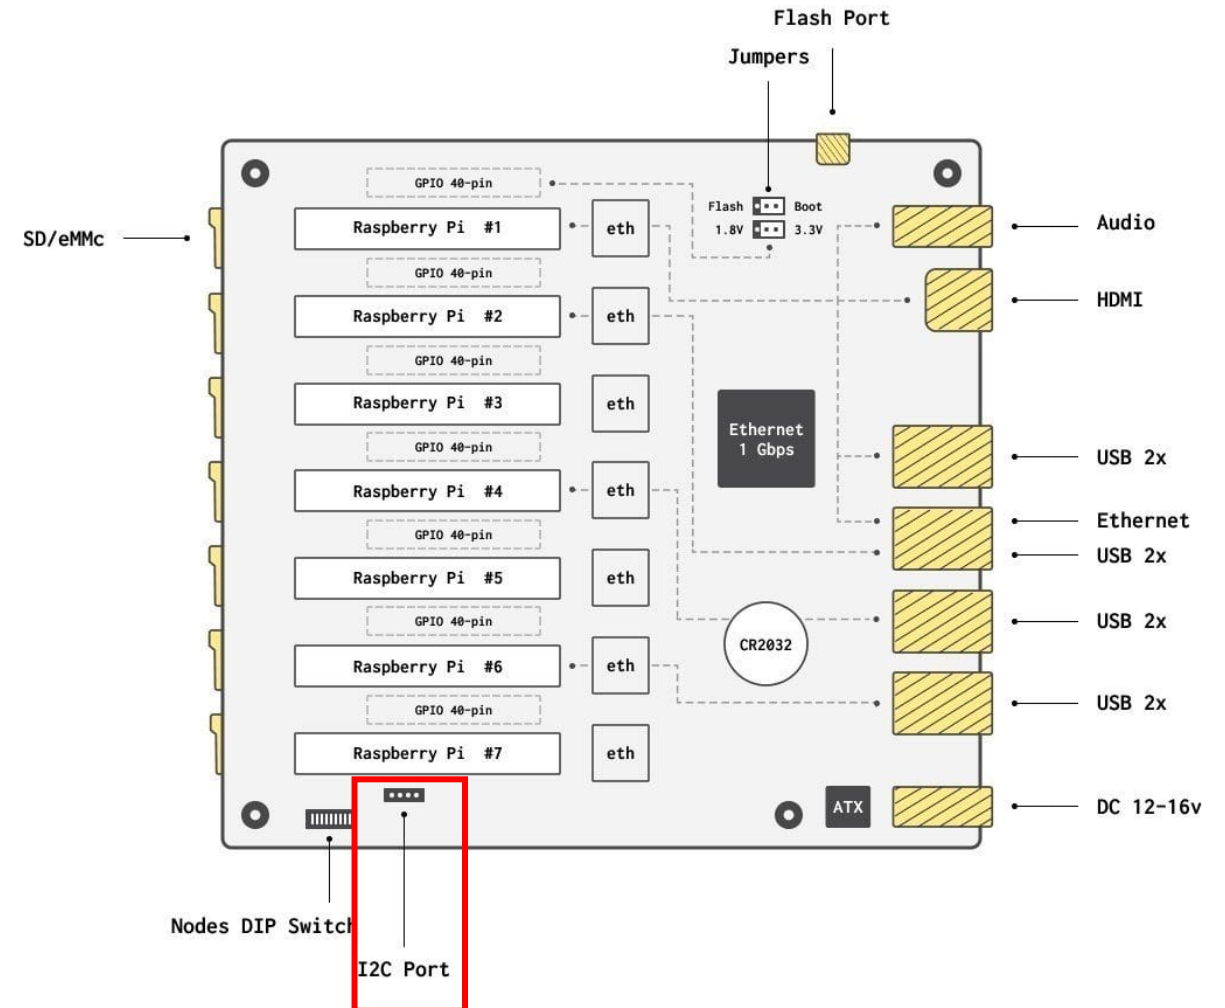


# Demo Infrastructure - turingpi v1



Kubernetes  
EDGE

- 7 RaspberryPi CM3+ modules
  - 1GB RAM
  - 4 CPU
- 8 USB ports
  - Map to only 4 devices
- I2C interface
  - Maps to all devices
  - Communication between devices
- GPIO
  - 40 pins for each device

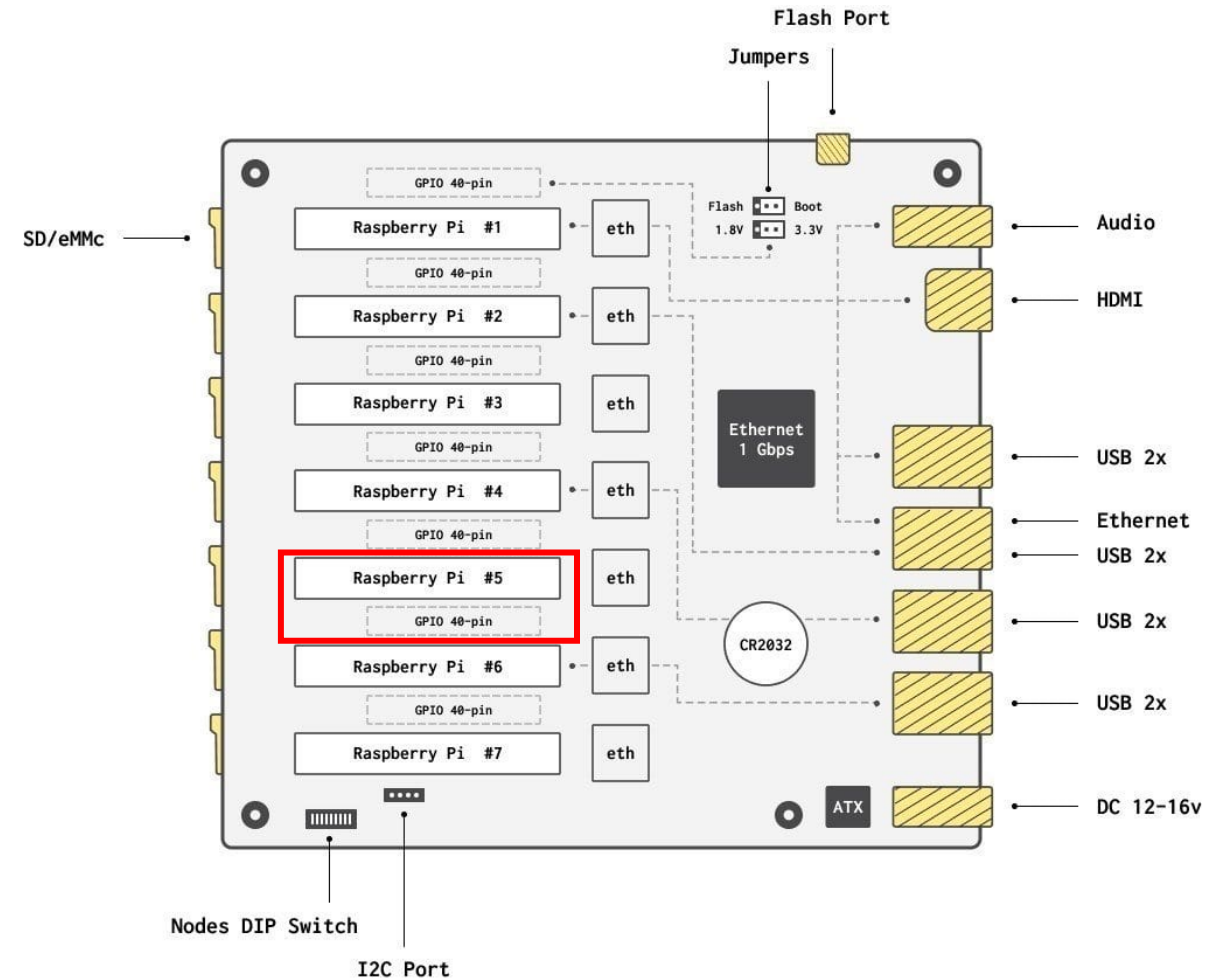


# Demo Infrastructure - turingpi v1



Kubernetes  
EDGE

- 7 RaspberryPi CM3+ modules
  - 1GB RAM
  - 4 CPU
- 8 USB ports
  - Map to only 4 devices
- I2C interface
  - Maps to all devices
  - Communication between devices
- **GPIO**
  - 40 pins for each device





**Kubernetes**  
**EDGE**

# Demo

# How does it work



**Kubernetes**  
**EDGE**

- Two Kubernetes Resources
  - NFD Master (Deployment)
  - NFD Worker (DaemonSet)

# How does it work



**Kubernetes**  
**EDGE**

- Operations team responsibilities
  - Know your devices
  - Create configuration defining known resources
  - NFD does the rest



# Install Node Feature Discovery



**Kubernetes**  
**EDGE**

- CRDs
  - K8s templates (yaml config)
  - Helm Chart
  - Operator
- NFD does what it does based on a configuration file
- Configmap is the key to the labeling

# Example ConfigMap



**Kubernetes**  
**EDGE**

```
#core:
#  labelWhiteList:
#  noPublish: false
#  sleepInterval: 60s
#  sources: [all]
#  klog:
#    addDirHeader: false
#    alsoLogToStderr: false
#    logBacktraceAt:
#    logToStderr: true
#    skipHeaders: false
#    stderrThreshold: 2
#    v: 0
#    vmodule:
#    logDir:
#    logFile:
#    logFileMaxSize: 1800
#    skipLogHeaders: false
```

```
sources:
#  cpu:
#    cpuid:
##      NOTE: whitelist has priority over
blacklist
#    attributeBlacklist:
#      - "BMI1"
#    attributeWhitelist:
#  kernel:
#    kconfigFile: "/path/to/kconfig"
#    configOpts:
#      - "NO_HZ"
#      - "X86"
#      - "DMI"
  usb:
    deviceClassWhitelist:
      - "03"
    deviceLabelFields:
      - "class"
      - "vendor"
      - "device"
```

# Practical NFD - Architecture Review



Kubernetes  
EDGE

- Data Center / Cloud

- dev



- test



- staging



- prod



# Practical NFD - Architecture Review



Kubernetes  
EDGE

- Data Center / Cloud

- dev



- test




- staging



- prod



- Edge - Its all production



# GitOps for Apps and NFD at the Edge



**Kubernetes**  
**EDGE**

- Rancher Fleet
- Argo CD
- Flux

# Example gitops workflow



Kubernetes  
EDGE

1. Install NFD via the gitops pipeline into every cluster
  2. Setup NFD configmap and commit to pipeline
  3. The system managing the clusters should automatically
    1. recognize changes in downstream cluster node labels and
    2. update the cluster labels accordingly
  4. Notify developers
    - of the node labels available
    - and what resource it maps to
- Repeat steps 2 - 4 on changes, hardware updates, etc.



1. Mutual TLS capable between nfd-master and nfd-worker
2. Risk: Allowing Local - user specific features
  1. will execute arbitrary files located in
  2. `/etc/kubernetes/node-feature-discovery/source.d` on the host
3. Limit features discovered by specifying exact sources
  1. Whilelist/Blacklist features
  2. `nfd-worker.conf`:
    1. set only specific sources
    2. e.g. sources: `[usb,custom,local]`



1. Label nodes in a cluster
2. NFD configuration is cluster scoped
  1. The edge will be thousands of single node or small clusters
  2. External tooling is needed to find groups of the same feature (gitops)
3. Feature discovery can be specified via configuration file
4. Limited features types
  1. Non kernel module features can be added using the *local* option
  2. Compare with Kubernetes device plugins





1. NFD allows drop in configmaps to another location
  1. Vendors can/should provide configs for their devices
  2. The NVIDIA GPU operator currently takes advantage of NFD
2. NFD is not “just” a labeler.
  1. It is a dynamic labeler
  2. It is a configurable labeler



1. Kernel module existence is not a guarantee of functionality
  1. Kernel modules can be loaded that have no physical counterpart
  2. K8s health checks may add value here
2. Hot plugging devices works but it is not immediate
  1. Kubernetes will need some cycles to label
  2. Re-labeling may also require redeployment of workloads



- <https://kubernetes-sigs.github.io/node-feature-discovery/v0.8/get-started/>
- <https://kubernetes-sigs.github.io/node-feature-discovery/v0.8/get-started/features.html>
- <https://www.usb.org/defined-class-codes>
- <https://k3s-io/k3s>
- <https://github.com/mak3r/nfd-demo>



**Kubernetes**  
**EDGE**

# Thank you

*Mark Abrams - Field Engineer and Edge Specialist*  
*SUSE*